

Bite Size Linux

By Julia Evans



♥ Table of contents ♥

unix permissions...4	sockets.....10	virtual memory...17
/proc.....5	unix domain sockets.....11	shared libraries...18
system calls.....6	processes.....12	copy on write19
signals.....7	threads.....13	page faults.....20
file descriptors..8	floating point....14	mmap.....21
pipes.....9	file buffering15	man page sections.....22
	memory allocation.....16	

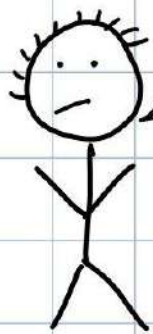
Linux fundamentals

by Julia Evans

pipes

mutexes

/proc



threads



file
descriptors

it's not so
bad! Let's
go learn the
main ideas!

signals



permissions



sockets

inodes

JULIA EVANS
@bork

unix permissions

drawings.jvns.ca

There are 3 things you can do to a file

↓
read write execute

ls -l file.txt shows you permissions
Here's how to interpret the output:

rw- rw- r-- bork staff
↑ ↑ ↑
bork (user) staff (group) ANYONE
can can
read & write read & write read

File permissions are 12 bits

setuid setgid
↓ ↓
000 110 110 100
sticky rwx rwx rwx

For the r/w/x bits:

1 means "allowed"

0 means "not allowed"

110 in binary is 6

So rw- r-- r--
= 110 100 100
= 6 4 4

chmod 644 file.txt
means change the permissions to:

rw- r-- r--
simple!

setuid affects executables

\$ls -l /bin/ping

rwS r-x r-x root root
↑
this means ping always runs as root

setgid does 3 different unrelated things for executables, directories, and regular files



an amazing directory: `/proc`

JULIA EVANS
@b0rk

Every process on Linux has a PID (process ID) like 42.

In `/proc/42`, there's a lot of VERY USEFUL information about process 42

`/proc/PID/cmdline`

command line arguments the process was started with

`/proc/PID/exe`

symlink to the process's binary. *magic*: works even if the binary has been deleted!

`/proc/PID/envIRON`

all of the process's environment variables

`/proc/PID/status`

Is the program running or asleep? How much memory is it using? And much more!

`/proc/PID/fd`

Directory with every file the process has open!

Run `$ls -l /proc/42/fd` to see the list of files for process 42.

These symlinks are also magic & you can use them to recover deleted files ♥

`/proc/PID/stack`

The kernel's current stack for the process. Useful if it's stuck in a system call

`/proc/PID/maps`

List of process's memory maps. Shared libraries, heap, anonymous maps, etc.

and `:more:`

Look at

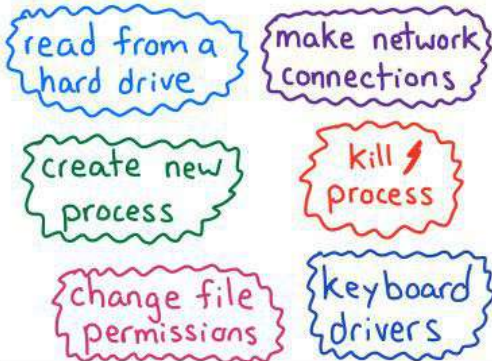
`man proc`

for more information!

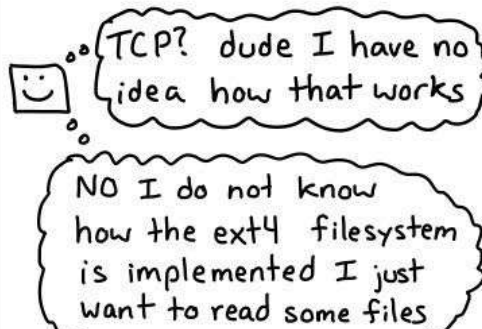
system calls

JULIA EVANS
@b0rk

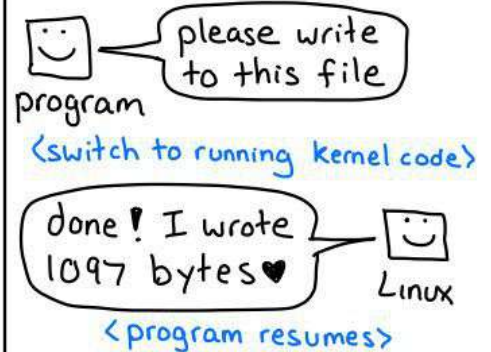
The Linux kernel has code to do a lot of things



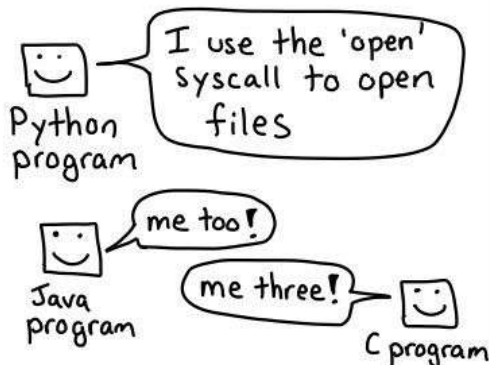
your program doesn't know how to do those things



programs ask Linux to do work for them using system calls

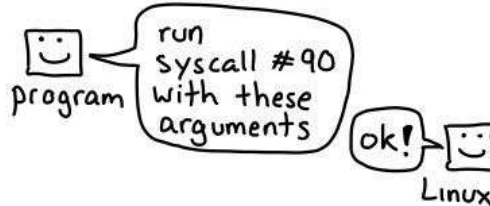


every program uses system calls



and every system call has a number (eg chmod is #90)

So what's actually going on when you change a file's permissions is



you can see which system calls a program is using with strace

```
$ strace ls /tmp
```

will show you every system call 'ls' uses! it's really fun!

⚠ strace is high overhead don't run it on your production database

JULIA EVANS
@b0rk

signals

drawings.jvns.ca

If you've ever used
⚡ kill ⚡
you've used signals



the Linux kernel sends
your process signals in
lots of situations



you can send signals
yourself with the **kill**
system call or command

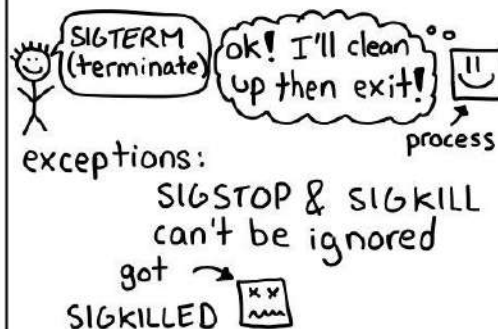
SIGINT ctrl-C } various
SIGTERM kill } levels of
SIGKILL kill -9 } "die"

SIGHUP kill -HUP
↑
often interpreted as
"reload config", eg by nginx

Every signal has a default
action, one of:

- ☺ ignore
- ☒ kill process
- ☒ ☹ kill process AND
make core dump file
- ⏸ stop process
- ⏪ resume process

Your program can set
custom handlers for
almost any signal



signals can be hard
to handle correctly since
they can happen at
ANY time



JULIA EVANS
@b0rk

file descriptors

Unix systems use integers to track open files



these integers are called **file descriptors**

lsf (list open files) will show you a process's open files

```
$ lsf -p 4242 ← PID we're interested in
```

FD	NAME
0	/dev/pts/tty1
1	/dev/pts/tty1
2	pipe:29174
3	/home/bork/awesome.txt
5	/tmp/

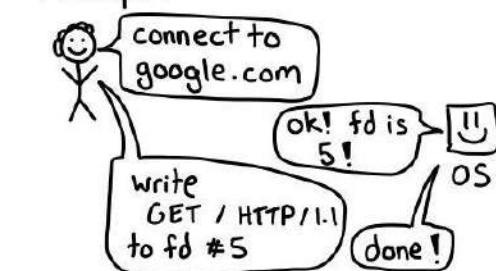
FD is for file descriptor

file descriptors can refer to:

- files on disk
- pipes
- sockets (network connections)
- terminals (like xterm)
- devices (your speaker! /dev/null!)
- LOTS MORE (eventfd, inotify, signalfd, epoll, etc etc)

not EVERYTHING on Unix is a file, but lots of things are

When you read or write to a file/pipe/network connection you do that using a file descriptor



Let's see how some simple Python code works under the hood:

```
Python:  
f = open("file.txt")  
f.readlines()
```

Behind the scenes:



(almost) every process has 3 standard FDs

- stdin → 0
- stdout → 1
- stderr → 2

"read from stdin"

means

"read from the file descriptor 0"

could be a pipe or file or terminal

JULIA EVANS
@bork

pipes

drawings.jvns.ca

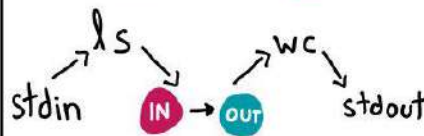
Sometimes you want to send the output of one process to the input of another

```
$ ls | wc -l
```

53
↖ 53 files!

a pipe is a pair of 2 magical file descriptors

IN and OUT



When `ls` does `write(IN, "hi")`

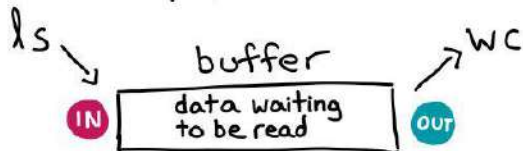
`wc` can read it!

`read(OUT)`

→ "hi"

Pipes are one-way. →
You can't write to OUT.

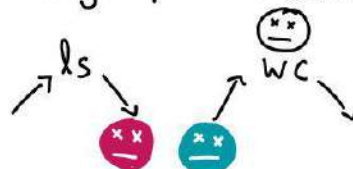
Linux creates a buffer for each pipe



If data gets written to the pipe faster than it's read, the buffer will fill up. IN [full buffer] OUT

When the buffer is full, writes to IN will block (wait) until the reader reads. This is normal & ok! 😊

what if your target process dies?



If `wc` dies, the pipe will close and `ls` will be sent SIGPIPE. By default SIGPIPE terminates your process.

named pipes

```
$ mkfifo my-pipe
```

This lets 2 unrelated processes communicate through a pipe!

`f=open("./my-pipe")`
`f.write("hi!\n")`

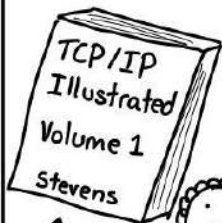
`f=open("./my-pipe")`
`f.readline() ← "hi!"`

JULIA EVANS
@b0rk

sockets

drawings.jvns.ca

networking protocols
are complicated



600 pages

what if I
just want to
download
a cat picture

Unix systems have
an API called the
"socket API" that
makes it easier to make
network connections
(Windows too! ☺)



Unix

you don't need to
know how TCP works,
I'll take care of it!

here's what getting
a cat picture with the
socket API looks like:

① Create a socket

`fd = socket(AF_INET, SOCK_STREAM, ...)`

② Connect to an IP/port
`connect(fd, 12.13.14.15:80)`

③ Make a request

`write(fd, "GET /cat.png HTTP/1.1 ...")`

④ Read the response

`cat-picture = read(fd ...)`

Every HTTP library uses
sockets under the hood

`$ curl awesome.com` → sockets
Python: `requests.get("yay.us")` → sockets



oh, cool, I could
write a HTTP
library too if I
wanted. * Neat!

* SO MANY edge cases
though! ☺

AF_INET?
What's that?

AF_INET means basically
"internet socket": it lets you
connect to other computers
on the internet using their
IP address.

The main alternative is
AF_UNIX ("unix domain socket")
for connecting to programs
on the same computer

3 kinds of internet
(AF_INET) sockets:

SOCK_STREAM = TCP

↑
curl uses this

SOCK_DGRAM = UDP

↑
dig (DNS) uses this

SOCK_RAW = just let me

↑
ping uses
this

send IP packets
I will implement
my own protocol

unix domain sockets

JULIA EVANS
@b0rk

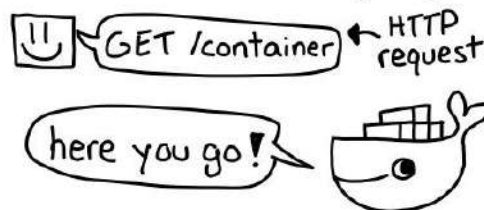
unix domain sockets
are files.

`$ file mysock.sock`
socket

the file's permissions
determine who can send
data to the socket

they let 2 programs
on the **same computer**
communicate.

Docker uses Unix domain
sockets, for example!



There are 2 kinds of
unix domain sockets:

stream like TCP! Lets
you send a
continuous
stream of bytes

datagram like UDP!
Send discrete
chunks of data



advantage 1

Lets you use file permissions
to restrict access to HTTP/
database services!

`chmod 600 secret.sock`

This is why Docker uses
a unix domain socket



advantage 2

UDP sockets aren't always
reliable (even on the same
computer).

unix domain datagram
sockets are reliable!
And won't reorder!



advantage 3

You can send a file
descriptor over a unix
domain socket.
Useful when handling untrusted
input files!



What's in a **process** ?

JULIA EVANS
@b0rk

PID

process #129
reporting for
duty!



USER and GROUP

who ran
you?

julia!



ENVIRONMENT VARIABLES

like PATH! you
can set them with:
`$ env A=val ./program`

SIGNAL HANDLERS



I ignore
SIGTERM!

I shut
down safely!



WORKING DIRECTORY

Relative paths (./blah)
are relative to the
working directory!
chdir changes it.

PARENT PID



PID 1 (init)
is everyone's
ancestor



PID 147



PID 129

COMMAND LINE ARGUMENTS

See them in
`/proc/PID/cmdline`

OPEN FILES

Every open file has
an offset.



I've read 8000
bytes of that one

MEMORY

heap! stack! ≡
shared libraries!
the program's binary!
mmaped files!

THREADS

sometimes one
sometimes LOTS

CAPABILITIES



I have
CAP_PTRACE

well I have
CAP_SYS_ADMIN



NAMESPACES



I'm in the host
network namespace

I have my own
namespace!



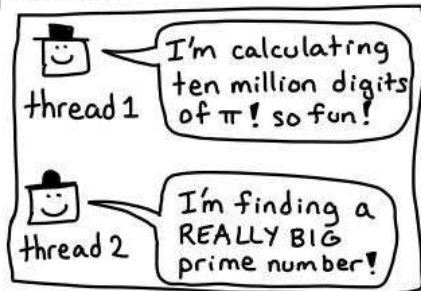
container
process

JULIA EVANS
@b0rk

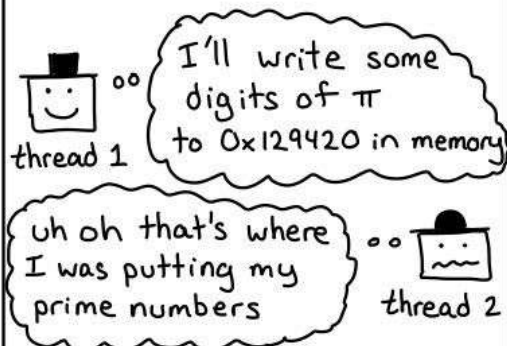
threads drawings.jvns.ca

Threads let a process do many different things at the same time

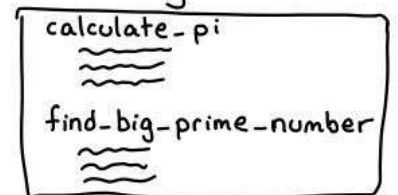
process:



threads in the same process share memory



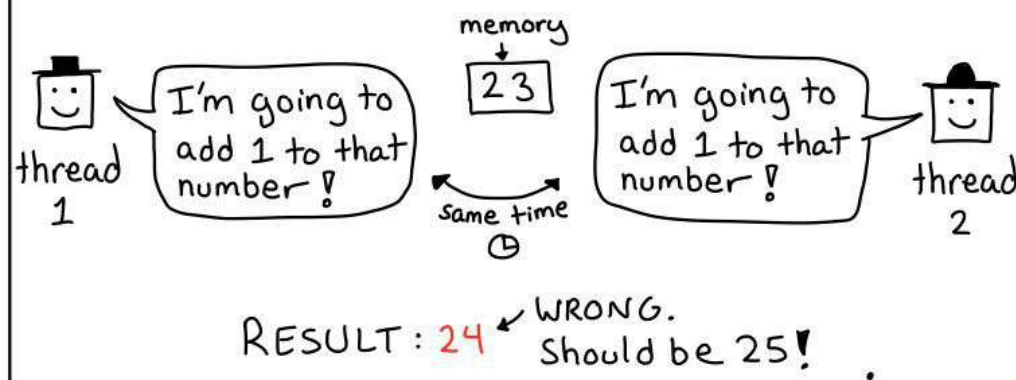
and they share code



but each thread has its own stack and they can be run by different CPUs at the same time



sharing memory can cause problems (race conditions!)



Why use threads instead of starting a new process?

→ a thread takes less time to create

→ sharing data between threads is very easy. But it's also easier to make mistakes with threads



floating point

JULIA EVANS
@bork

a double is 64 bits

sign exponent fraction
↓ ↓ ↓
10011011 10011011 10011011 10011011
10011011 10011011 10011011 10011011

$$\pm 2^{e-1023} \times 1.\text{frac}$$

That means there are
 2^{64} doubles

The biggest one is about
 2^{1023}

weird double arithmetic

$$2^{52} + 0.2 = 2^{52}$$

← (the next number after
 2^{52} is $2^{52} + 1$)

$$1 + \frac{1}{2^{54}} = 1$$

← (the next number after
1 is $1 + \frac{1}{2^{52}}$)

$$2^{2000} = \text{infinity}$$

← infinity is a double

$$\text{infinity} - \text{infinity} = \text{nan}$$

← nan = "not a number"

doubles get farther
apart as they get bigger

between 2^n and 2^{n+1} there
are always 2^{52} doubles,
evenly spaced

that means the next double
after 2^{60} is $2^{60} + 64 \leftarrow \frac{2^{60}}{2^{52}}$

Javascript only has
doubles (no integers!)

$$> 2^{**}53$$

9007199254740992

$$> 2^{**}53 + 1$$

9007199254740992

↑
same number! uh oh!



doubles are scary
and their arithmetic
is weird

they're very logical!
just understand how
they work and don't
use integers over 2^{53}
in Javascript ♥



file buffering

???

I printed some text but it didn't appear on the screen. why??

time to learn about **flushing!**

On Linux you write to files & terminals with a system call called

♥ **write** ♥

please write "I ♥ cats" to file #1 (stdout)

okay!

Linux

I/O libraries don't always call **write** when you print

```
printf("I ♥ cats");
```

°° I'll wait for a newline before actually writing

This is called **buffering** and it helps save on syscalls

3 kinds of buffering (defaults vary by library)

- ① None. This is the default for **stderr**
- ② Line buffering. (write after newline). The default for **terminals**.
- ③ "full" buffering. (write in big chunks). The default for **files** and **pipes**.

flushing 

To force your IO library to write everything it has in its buffer right now, call **flush!**

I'll call **write** right away!!

stdio

when it's useful to flush

- when writing an interactive prompt!

Python example:

```
print("password: ", flush=True)
```

- when you're writing to a pipe / socket

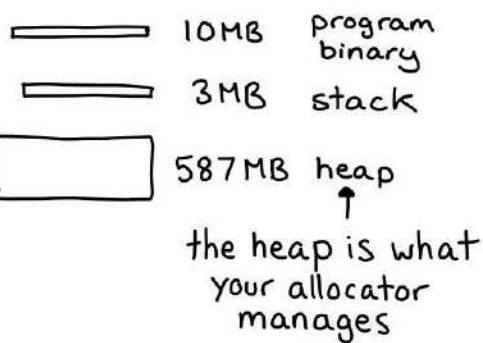
no seriously, actually write to that pipe please

program

memory allocation

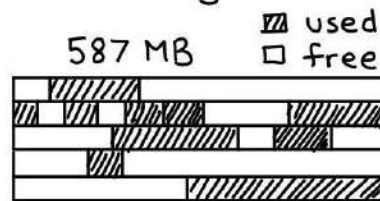
JULIA EVANS
@bork

your program has memory

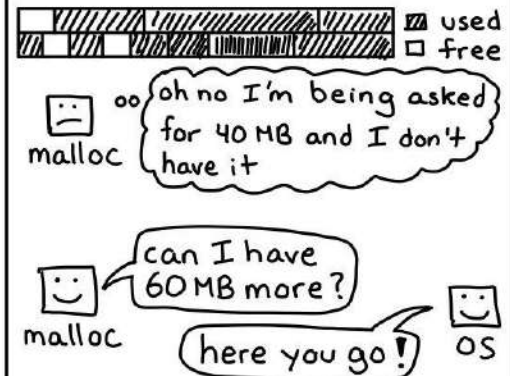


your memory allocator (malloc) is responsible for 2 things.

THING 1: keep track of what memory is used/free



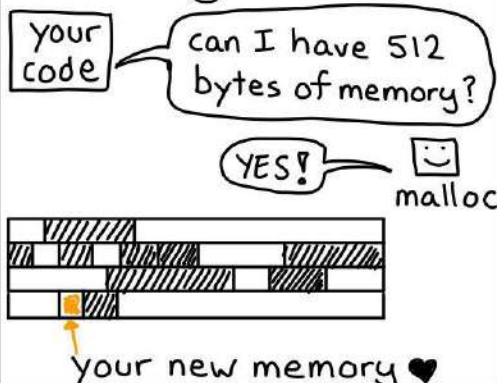
THING 2: Ask the OS for more memory!



your memory allocator's interface

malloc (size_t size)
allocate **size** bytes of memory & return a pointer to it
free (void* pointer)
mark the memory as unused (and maybe give back to the OS)
realloc (void* pointer, size_t size)
ask for more/less memory for **pointer**
calloc (size_t members, size_t size)
allocate array + initialize to 0

malloc tries to fill in unused space when you ask for memory



malloc isn't magic!
it's just a function!

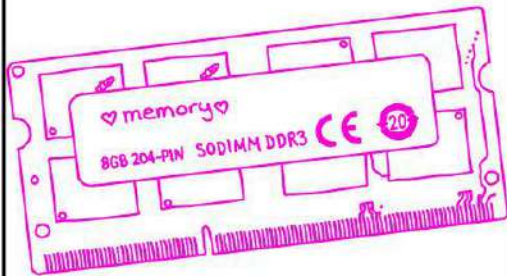
you can always:

- use a different malloc library like **jemalloc** or **tc malloc** (easy!)
- implement your own malloc (harder)

virtual memory

JULIA EVANS
@b0rk

your computer has
physical memory



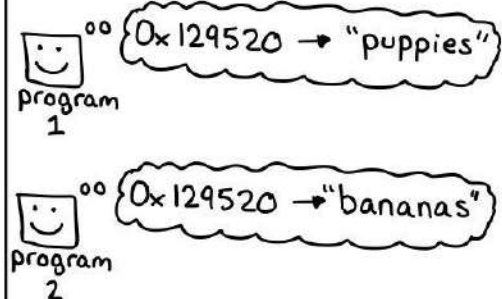
physical memory has
addresses

0-8GB

but when your program
references an address
like 0x5c69a2a2

that's not a physical
memory address!
It's a **virtual** address

every program has its
own virtual address space

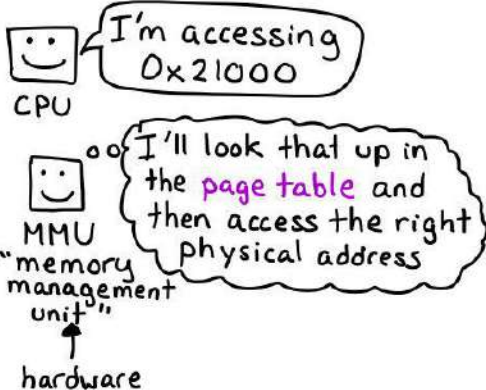


Linux keeps a mapping from
virtual memory pages to
physical memory pages called
the "**page table**"

a "page" is a 4kb* or
chunk of memory sometimes
bigger

PID	virtual addr	physical addr
1971	0x20000	0x192000
2310	0x20000	0x228000
2310	0x21000	0x9788000

when your program
accesses a virtual address



every time you switch
which process is running,
Linux needs to switch
the page table



shared libraries

JULIA EVANS
@b0rk

Most programs on Linux
use a bunch of C libraries

some popular libraries:

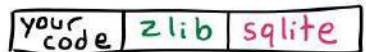
openssl (for SSL!) sqlite (embedded db!)

libpcre (regular expressions!) zlib (gzip!)


libstdc++ (C++ standard library!)

There are 2 ways
to use any library

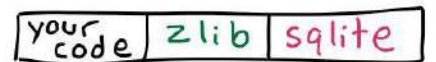
- ① Link it into your binary


big binary with lots of things!

- ② Use separate shared libraries

 ← all different files

Programs like this:



are called "statically linked"

and programs like this:



are called "dynamically linked"

how can I tell what
shared libraries a
program is using?

ldd!!

```
$ ldd /usr/bin/curl
libz.so.1 => /lib/x86_64..
libresolv.so.2 => ...
libc.so.6 => ...
+ 34 more ☺
```

I got a "library not
found" error when running
my binary?!

If you know where the
library is, try setting
the **LD_LIBRARY_PATH**
environment variable

☺ ∞
dynamic linker **LD_LIBRARY_PATH**
tells me where to look!

Where the dynamic
linker looks

- ① DT-RPATH in your executable
- ② LD_LIBRARY_PATH
- ③ DT-RUNPATH in executable
- ④ /etc/ld.so.cache
(run ldconfig -p to see contents)
- ⑤ /lib, /usr/lib

copy on write

JULIA EVANS
@b0rk

On Linux, you start new processes using the `fork()` or `clone()` system call

calling fork gives you a child process that's a copy of you



parent



child

the cloned process has EXACTLY the same memory

- same heap
- same stack
- same memory maps

if the parent has 3GB of memory, the child will too

copying all that memory every time we fork would be **slow** and a **waste of RAM**



often processes call `exec` right after `fork` which means they don't use the parent process's memory basically at all!

so Linux lets them share physical RAM and only copies the memory when one of them tries to **write**.



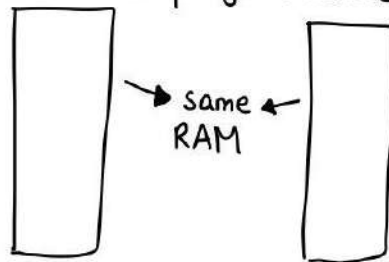
I'd like to change that memory

ok I'll make you your own copy!



Linux

Linux does this by giving both the processes identical page tables



but marks every page as **read only**

When a process tries to write to a shared memory address

- ① there's a **page fault**
- ② Linux makes a copy of the page & updates the page table
- ③ the process continues, blissfully ignorant



It's just like I have my own copy

page faults

JULIA EVANS
@b0rk

every Linux process has a page table

★ page table ★

virtual memory address	physical memory address
0x19728000	0x1422000
0x19724000	0x1423000
0x1524000	not in memory
0x1844000	0x4a000 read only

some pages are marked as either

- ★ read only
 - ★ not resident in memory
- when you try to access a page that's marked "not in memory", that triggers a **! page fault!**

What happens during a page fault?

- the MMU sends an interrupt
- your program stops running
- Linux kernel code to handle the page fault runs

Linux 😊 "I'll fix the problem and let your program keep running"

"not in memory" usually means the data is on disk!

virtual memory



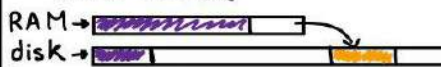
Having some virtual memory that is actually on disk is how **swap** and **mmap** work

how swap works

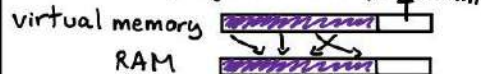
① run out of RAM



② Linux saves some RAM data to disk



③ mark those pages as "not resident in memory" in the page table



④ When a program tries to access the memory there's a **! page fault!**

⑤ 😊 Linux "time to move some data back to RAM!"



⑥ if this happens a lot your program gets **VERY SLOW**

😞 "I'm always waiting for data to be moved in & out of RAM"

JULIA EVANS
@b0rk

mmap

drawings.jvns.ca

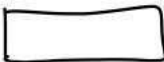
What's mmap for?

I want to work with
a VERY LARGE FILE
but it won't fit
in memory

you could
try mmap!
(mmap = "memory map")

load files lazily
with mmap


When you mmap a file, it
gets mapped into your
program's memory

2TB file  2TB of
virtual memory
but nothing is ACTUALLY
read into RAM until you
try to access the memory
(how it works: page faults!)

how to mmap
in Python

```
import mmap
f = open("HUGE.txt")
mm = mmap.mmap(f.fileno(), 0)
    ↗ this won't read the
    file from disk!
    Finishes ~instantly.
print(mm[-1000:])
    ↑
    this will read only
    the last 1000 bytes!
```


sharing big files
with mmap

 we all want to
- read the same file!
no problem! mmap

Even if 10 processes
mmap a file, it will only
be read into memory
♥ once ♥

dynamic linking
uses mmap

 program I need to
use libc.so.6
↗ (standard library)

you too eh? no problem
I always mmap, so
that file is probably
loaded into memory
already  dynamic
linker

anonymous memory maps

- not from a file
(memory set to 0 by default)
- with MAP_SHARED, you can
use them to share memory
with a subprocess!

man pages = awesome

JULIA EVANS
@b0rk

man pages are split up
into 8 sections

① ② ③ ④ ⑤ ⑥ ⑦ ⑧

\$ man 2 read

means "get me the man page
for read from section 2"

There's both

- a program called "read"
 - and a system call called "read"
- so

\$ man 1 read

gives you a different man page from

\$ man 2 read

If you don't specify a section, man will
look through all the sections & show
the first one it finds

man page sections

① programs

\$ man grep
\$ man ls

② system calls

\$ man sendfile
\$ man ptrace

③ C functions

\$ man printf
\$ man fopen

④ devices

\$ man null
for /dev/null docs

⑤ file formats

\$ man sudoers
for /etc/sudoers
\$ man proc
files in /proc!

⑥ games

not super useful.
\$ man sl
is funny if you have
sl though.

⑦ miscellaneous

explains concepts!

\$ man 7 pipe
\$ man 7 symlink

⑧ sysadmin programs

\$ man apt
\$ man chroot